




Your Thesis Title Here



Rafael Tanzer

Cyber-Physical-Systems, ..., ...  
TU Vienna

Supervisor: Michele Chiaria, Simone , Francesco Pontigia and Ezio Bartocci

May 2025

This thesis is submitted in partial fulfillment of the requirements for the degree of *Bachelor of Science*.

## Abstract

# Contents

<b>I</b>	<b>Foundations</b>	<b>3</b>
<b>1</b>	<b>Introduction</b>	<b>4</b>
1.1	Motivation and Problem Statement . . . . .	4
1.2	Objectives and Scope . . . . .	4
<b>2</b>	<b>Background on Languages &amp; Grammars</b>	<b>5</b>
2.1	Formal Language Theory . . . . .	5
2.1.1	Backus–Naur Form (BNF) and Variants . . . . .	6
2.1.2	Grammar Formalisms and Parsing Models . . . . .	6
2.2	Language Structure and Paradigms . . . . .	6
2.2.1	Abstract vs. Concrete Syntax . . . . .	6
2.2.2	Imperative vs. Declarative Languages . . . . .	7
2.3	Domain-Specific Languages (DSLs) . . . . .	7
2.3.1	Imperative DSLs . . . . .	8
<b>3</b>	<b>Probabilistic Programming &amp; Cyber-Physical Systems</b>	<b>9</b>
3.1	What Is Probabilistic Programming? . . . . .	9
3.2	Inference and Sampling Methods . . . . .	9
3.3	Cyber-Physical Systems: Definition and Examples . . . . .	10
3.4	Motivation for PPLs in CPS Contexts . . . . .	11
<b>4</b>	<b>The <i>Miniprob</i> DSL</b>	<b>13</b>
4.1	Domain and Purpose . . . . .	13
4.2	Syntax Overview . . . . .	13
4.3	Semantics and Example Models . . . . .	13
<b>5</b>	<b>Language Servers, Parsers &amp; Type Checking</b>	<b>14</b>
5.1	Language Server Protocol (LSP) Basics . . . . .	14
5.2	Parser and Syntax Checking . . . . .	14
5.3	Type Checking Functionality . . . . .	14
5.4	Overview of Langium . . . . .	14
<b>II</b>	<b>Design &amp; Implementation</b>	<b>15</b>
<b>6</b>	<b>Requirements &amp; Technology Survey</b>	<b>16</b>
6.1	Functional Requirements . . . . .	16
6.2	Non-functional Requirements . . . . .	16
6.3	Alternative Technologies . . . . .	16
6.4	Justification for Langium & VS Code . . . . .	16

7	Architectural Design	17
7.1	High-Level Architecture Diagram	17
7.2	Module Decomposition	17
7.3	Data Flow and Control Flow	17
8	Implementation Details	18
8.1	Langium Grammar Definition for <i>Miniprob</i>	18
8.2	Semantic Checks and Type System	18
8.3	VS Code Extension Points	18
9	Testing & Validation	19
9.1	Unit Tests	19
9.2	Integration Tests	19
9.3	Case Studies & Examples	19
III	Evaluation & Discussion	20
10	Performance & Usability	21
10.1	Parsing Speed	21
10.2	Editor Responsiveness	21
10.3	User Feedback	21
II	Discussion	22
II.1	Meeting the Requirements	22
A	Expanded Code Listings	25
B	Test Data	26

Part I

Foundations

# Introduction

## 1.1 Motivation and Problem Statement

CPS team - waht do they need it for code growing user base - interconnecticty os , other qualities speaking for the usage of vs code - already in possession of Haskell Checker but

## 1.2 Objectives and Scope

The main goal of the thesis was to create a *Language Extension* ?NAME? for *MiniProb*. Language extensions are a big part of VS Code’s extension ecosystem and enable the editor to utilize custom language tooling. Such extensions support the user during the implementation process while writing code, by providing language assistance like syntax validation or code completion. While this implementation targets VS Code, the underlying logic and functionality are not limited to it — through the use of the *Language Server Protocol (LSP)*, the developed tooling can be reused across various editors and IDEs that support the protocol. The core functionality of these extensions stems from *parsers* which, against a formal language definition, convert written text into abstracted parts, validating the code in the process. These parts can then be used to extend the capabilities of the tooling to encompass referential validation, type checking, code completion, diagnostics, and other language services that enhance the development experience.

In order for parsers to function accordingly, they require formal language definitions or models. These definitions range from various types of grammars to abstract machines such as automata, which together provide the structural and syntactic rules necessary for correct interpretation and validation of source code. Based on these definitions, parsers are able to systematically identify the hierarchical structure of a program by recognizing patterns, token sequences, and nested constructs, ensuring that the code adheres to the expected form before any further processing or analysis occurs. In this thesis, a generative context-free grammar is used to formally describe *MiniProb*, a domain-specific language (DSL) designed for authoring *POMC* files. Based on this grammar, a parser - serving as the foundation - enables the implementation of a fully functional language extension to aid developers writing *POMC* files.

Ultimately, the resulting extension, titled ?NAME?, is intended to provide comprehensive language support for *MiniProc*. This includes syntactic analysis through syntax highlighting and validation, accurate resolution of symbol references and code completion as well as the implementation of type checking mechanisms to ensure semantic correctness during development. //maybe basic code completion for expressions

An appropriate set of regression tests was developed to ensure the continued correctness and stability of the language tooling. These tests include parsing tests, which verify that valid input is correctly recognized and structured according to the grammar; validation tests, which ensure that semantic rules are properly enforced and errors are accurately reported; and linking tests, which check that references between symbols or declarations are correctly resolved across different parts of a program. Together, these test categories help maintain the integrity of the parser and language services as the implementation evolves.

Lastly, this thesis includes an evaluation of the newly implemented parser, focusing on its correctness and performance in practical usage scenarios. The evaluation is based on metrics collected from representative input samples, measuring factors such as parsing speed, memory usage, and error handling. Where relevant, comparisons are also drawn against the existing *Haskell*-based parser for *.pomc* files, offering a point of reference to assess improvements or trade-offs introduced by the new implementation.

# Background on Languages & Grammars

## 2.1 Formal Language Theory

Formal language theory deals with the study of languages – sets of strings constructed from alphabets – and the formal grammars that determine and generate them. In contrast to natural languages, which have evolved over centuries under the influence of diverse cultural, historical, and environmental factors, formal languages do not inherently relate to any perceived constructs of our environments and are generally not intuitively understood. Additionally, formal languages do not share the rich evolutionary progression — a lengthy process of gradual adaptations and refinements spanning generations — of their natural counterparts. Instead, they employ sets of axiomatic *production rules* that describe each language individually. The field of formal language theory sprung from linguist Noam Chomsky’s attempts during the 1950s to definitively characterize the structure of natural languages using mathematical rules. [Jia+09] This analytical approach led to development of the *Chomsky hierarchy*, which proved to be a vital theoretical foundation for later discoveries and applications, as it was found that all information (photos, videos, numbers, axioms) can be represented as finite strings.

The *Chomsky Hierarchy* is a hierarchical classification of formal grammars, that labels them to four groups. The grammars are ranked based on the individual *expressive power* of the languages they produce, with each class including the less expressive ones. The whole hierarchy consists of four groups of grammars (and corresponding language classes), that are identified by inspecting the production rules, which get progressively less restrictive.

Type-3 grammars, or *regular* grammars, generate exactly the class of regular languages. Their productions are restricted to the two equivalent styles:

$$\text{right-}reg.: \quad A \rightarrow aB \quad \text{or} \quad A \rightarrow a \quad \text{and} \quad \text{left-}reg.: \quad A \rightarrow Ba \quad \text{or} \quad A \rightarrow a,$$

where  $A, B$  are nonterminals and  $a$  is a terminal. Regular expressions provide an alternative, declarative notation for these languages — each expression can be mechanically transformed into a regular grammar, and vice versa. More broadly, every grammar in the Chomsky hierarchy admits an equivalent acceptor automaton; in the case of Type-3 grammars, this correspondence yields deterministic or nondeterministic finite automata. Leveraging these conversions makes regular grammars invaluable in practice (for example, in lexical analysis, pattern matching, and protocol verification), even though their expressive power is the most limited. Type-2 grammars, or *context-free* grammars (CFGs), produce all context-free languages. Their rules take the general form

$$A \rightarrow \alpha$$

where  $A$  is a single nonterminal and  $\alpha$  is any string of terminals and nonterminals. This additional flexibility captures nested, hierarchical structures—like balanced parentheses or most programming-language syntaxes—that regular grammars cannot. Each CFG is accepted by a corresponding pushdown automaton: the grammar expansions map naturally onto push and pop operations on the PDA’s stack, making the grammar-automaton equivalence at this level another cornerstone of formal language theory.

for choosing between replacing or recalibrating components—but measured data often reflect intertwined effects and closed-loop corrections. titcontext-sensitive and *recursively enumerable* languages respectively, and are the most expressive of all with Type-0 grammars placing absolutely no constraints on the production, making them only acceptable by Turing-Machine.

### 2.1.1 Backus–Naur Form (BNF) and Variants

Even though production rules, alphabet specifications, and automaton definitions suffice to unambiguously define the set of valid strings/sentences in a language, their notation tends to be opaque and cumbersome in practice, prompting interest in forming intuitively comprehensible grammar definitions. The Backus-Naur Form (BNF), created by John Backus with contributions by Peter Naur and released in their Algol-60 report. [Bac+60], was a pivotal introduction furnishing a concise, human-readable syntax for language generation. By expressing each rule as

$$\langle \text{nonterminal} \rangle ::= \text{expansion}_1 \mid \text{expansion}_2 \mid \dots$$

BNF allows both sequencing and alternation of multiple expansions, enabling language designers to articulate complex, nested structures without exposing the underlying automaton states or transition tables. This declarative approach not only supports rigorous specification but also facilitates mechanical parsing, thereby advancing the practice of compiler development and language tooling.

While BNF provides a clear, formal way to describe language syntax, it suffers from verbosity and redundancy — pure BNF grammars often become bloated when encoding common patterns like optionals or repetitions requiring auxiliary nonterminals for a sufficient description and, over time, has been extended and modified for various use-cases spawning a new family of grammar notations. The **Extended Backus-Naur Form** extends BNF by adding more expressive meta-syntax, introducing operators similar to *Regular Expressions* allowing the use of optional or repeatable expressions, commonly indicated by brackets but not limited to '[...]' and '{...}', and the use of comments. Such extensions make grammars more compact and easier to read without changing the fundamental class of languages they describe.[] Multiple distinct instances of EBNF's have been development, all with minor syntactic differences, yet there really is no on-for-all EBNF used in practice, although a standardized ISO/IEC 14977 version exists. [Jinb; Jina]

Within the scope of my thesis, *?NAME?*, Extended Backus–Naur Form (EBNF) is particularly significant, since Langium employs its own custom EBNF dialect to define the grammars underpinning its language-assistance features. 8

### 2.1.2 Grammar Formalisms and Parsing Models

..programming languages

## 2.2 Language Structure and Paradigms

### 2.2.1 Abstract vs. Concrete Syntax

Depednming on subsection Grammar Formalisms, mentino the importance of grammars in connectino with programming langs.

The terms *abstract* and *concrete* syntax are primarily encountered in the context of programming language design. While the two concepts follow the common relationship of abstract and concrete instances — where one embodies a generalized, meta-level blueprint, while the other encapsulates the concrete, instance-level details — and the term "syntax" broadly applies to all languages, the idea of pinning meta-information onto language constructs, somewhat implies further calculations based on the language itself.

When designing programming languages, the abstract syntax is of particular interest, as its streamlined view of the language is ideally suited for developing validation constraints or type systems and for general interpretation. This is done most commonly by encoding the syntax into a **Abstract Syntax Tree (AST)**, a tree-like structure containing nodes for each high-level construct (such as expressions, declarations, or statements) connected according to their logical hierarchy, while deliberately omitting lexical details to yield a canonical. [SK95] The **Concrete Syntax Tree (CST)** is the specialized counterpart to the AST and consists of a full parse-tree, preserving every terminal and nonterminal token(keywords, delimiters, etc.) and mirrors each grammar production in its node structure, providing the complete syntactic context possibly needed for precise error reporting and tooling. [Aho+o6]



### 2.2.2 Imperative vs. Declarative Languages

Languages used for further computation can be analyzed from multiple perspectives and classified with different qualities in mind. [Van12] For this thesis we focus on slotting languages into two overarching paradigms, **imperative** and **declarative** languages, which differentiate the by inspecting way of reaching the desired outcome. **Declarative** languages describe *what* the computation should accomplish by adhering to sets of expressions, constraints or logical statements. [WGMo8] These formulas set the rules and goals characterizing the desired result or relationship of input and output values. The actual control flow and low-level decisions of the execution is left to the specific language implementations and runtime engines.

Generally, because of the vague wording used, the paradigms themselves are not always completely and accurately describing the behaviour of a programming language, possibly leading to different concepts overlapping. This difference is especially clear in modern languages, which often implement functionality by borrowing constructs from the opposite paradigm — for example, imperative C# incorporates declarative features through its LINQ-Library, while Haskell, functional by design, supports imperative-style programming. [netHaskellClaims]

Imperative languages, in contrast to declarative ones, describe *how* a computation will accomplish the desired end-state by following explicit sequences of commands. These commands consist of statements that actively manipulate the computations state, which is typically managed through assignments to variables and taking charge of the control flow with provided language constructs (loops, conditionals, calls, procedures, etc.). Consequently, correct execution is no longer responsibility of the underlying system, but the implementation and therefore the programmer itself. The imperative paradigm reflect von Neuman machine — accessing memory and hardware directly. This fined grained control enables the use of low-level optimization techniques [lowLevelOpt]. As programmers instruct the machine how to do something, certain performance characteristics and resource constraints can be impelled depending on the context of the computation. Such techniques encompass memory and cache management, *Instruction-Set* optimization or compiler optimization.

Verbose is missing?

However, while this explicit control provides developers with great functional freedom, it also places a greater burden on them to manage every detail of execution. Because the programmer must specify each step — from how data is stored in memory to the order in which operations occur — imperative code can become verbose and cluttered obscuring the high-level intent. Moreover, mutable state and side effects introduce the risk of subtle bugs, especially in concurrent or parallel settings where unintended interactions between state changes can lead to race conditions. Also, formal reasoning and verification are more complex, since proving correctness requires tracking every state transition rather than relying on *referentially transparent* expressions.

Typical use cases for imperative programming include algorithmic or process-oriented problems where an explicit series of steps is natural. Low-level system programming, embedded development, and performance-sensitive routines frequently rely on imperative constructs to manage hardware resources directly. In scenarios that demand fine-grained stateful interactions — such as updating user interfaces incrementally, reading and writing files in a specific order, or controlling physical systems via commands — imperative languages shine by giving developers precise command over execution. Simulation engines, game logic, and scripting or “recipe”-style automation languages also favor imperatively describing each step to achieve the desired outcome.

## 2.3 Domain-Specific Languages (DSLs)

for choosing between replacing or recalibrating components—but measured data often reflect intertwined effects and closed-loop corrections. to also domain experts. Moreover, DSLs declutter code, rather than describing a subproblem of the domain with multiple lines of code, fewer are needed because constructs inherently carry domain specific information, also contributing to broader comprehensibility for those familiar with the domain. DSLs solutions are also easier to formally verify, as they are restricted to a limited context, than GLPs for which proving correctness is extremely hard.

However, DSLs can suffer from limited scope and flexibility even within their target domain, forcing developers to fall back on general-purpose languages when they need functionality outside the DSL’s narrow focus. Furthermore, if the domain itself is complex or rapidly evolving, creating and maintaining a DSL may still de-

mand deep domain expertise, because the abstractions must accurately capture all necessary domain rules and nuances.

### 2.3.1 Imperative DSLs

Whichever paradigm a DSL follows is mostly guided by the specifics of the domain to describe. Even though DSLs usually adhere to a declarative approach [DKV00], which follows from a major motivation: wanting to describe *what* is expected of the domain in high-level abstraction, a domain may emit qualities that invoke more imperative thinking. If the nature of a domain is inherently composed of sequential actions, consists of procedures or stateful operations, the imperative paradigm is more appropriate. Imperative DSLs require the user to script solutions in step-by-step fashion using domain-specific commands and statements. Essentially, the user embeds domain concepts as a sequential programming model, where the control flow and state transitions of the domain computations have to be managed.

In domains where fine-grained control and explicit ordering mat for choosing between replacing or recalibrating components—but measured data often reflect intertwined effects and closed-loop corrections. reating a declarative one: typically, an interpreters process each command in sequence, or a translator converts DSL instructions into a host general-purpose language.

Mybe just leave the below out

Imperative DSLs shine in scripting and macro contexts, where users sequence domain-specific commands to drive behavior, examples include Excel macros that automate spreadsheet tasks or scripts within a game engine that dictate character actions and game logic. They are also well suited to testing frameworks and scenario specifications, where user interfaces describe interactions by specifying a timely sequence of user inputs. In robotics and industrial automation, imperative DSLs allow practitioners to program custom action sequences — such as moving a robotic arm to pick up an object, then placing it precisely — where each step must follow a strict order to ensure safety and accuracy. More broadly, any process-control domain that relies on precise sequencing and stateful interactions — be it a complex simulation workflow, an interactive operation pipeline, or automated machinery — benefits from the clarity and control that an impera for choosing between replacing or recalibrating components—but measured data often reflect intertwined effects and closed-loop corrections. tive DSL provides.

# Probabilistic Programming & Cyber-Physical Systems

## 3.1 What Is Probabilistic Programming?

Probabilistic programming is a paradigm that aims to perform statistical analysis using tools yielded by computer science. [Mee+21] Unifying general-purpose programming and probabilistic modelling enables the presentation of statistical models as a program. Defined by underlying program code, the model is embodied by the relations between variables and calculations. The established capabilities of programming languages are often augmented with probabilistic functionality, in ... creating probabilistic programming languages (PPL), [Goo13] that are closely associated with this paradigm. The enriched syntax enables the creation of more compact models and fosters conciser abstractions of the underlying probabilistic structure, usually providing constructs that allow declaring random variables and conditioning of observed data, typical for probabilistic modelling. These constructs form the foundation for specifying core elements of statistical models. *Priors* are pre-known and usually drawn from distributions reflecting the initial belief over what value *latent variables*/model parameters might take on. *Likelihoods* are also known beforehand representing how data is assumed to be generated given the latent variable for which the *posterior* is to be inferred. Together, these components form the backbone of a probabilistic model and provide the information required to compute the posterior distribution, which is grounded in **Bayesian theory**. [Mee+21]

*Bayes' Theorem* provides the foundational framework for probabilistic reasoning under uncertainty. It describes how to update prior beliefs about unknown quantities - known as latent variables - based on new evidence, resulting in a posterior distribution. Formally, it states that the posterior is proportional to the product of the prior and the likelihood:

$$P(\theta \mid x) \propto P(x \mid \theta) \cdot P(\theta)$$

In probabilistic programming, this principle is applied through models that define prior distributions over latent variables and likelihood functions for observed data. The probabilistic programming language then applies Bayes' Theorem, typically via automated inference algorithms, to compute or approximate the posterior. This makes it possible to express complex probabilistic models declaratively, while delegating the computation of posterior beliefs to the underlying inference engine.

??INCLUDE GRAPHICAL MODESL??

## 3.2 Inference and Sampling Methods

Inference is the process by which probabilistic models are "executed". Given a model and observed data, the task is to compute the posterior distributions of latent variables. Because probabilistic models often include uncertainty and hidden variables, inference is the process that turns the model into concrete answers. This process is abstracted and automated, but remains computationally nontrivial. Standard approaches to inference are generally grouped into three main categories- *exact inference*, *sampling-based methods*, and *variational inference*. [Biso6; BKM17; GS14].

### 1. Exact Inference

Exact inference computes the posterior distribution analytically or through complete enumeration. This is possible only in models with limited complexity where the entire state space can be traversed. Although conceptu-

ally straightforward, exact inference quickly becomes infeasible as model size increases, due to the combinatorial explosion of possible states.

## 2. Sampling-Based Inference

Sampling-based methods approximate the posterior in probabilistic programming,  $t$

Variational inference formulates inference as an optimization problem. Instead of sampling, it produces a family of approximate distributions and seeks the member that is closest to the true posterior. This involves defining an objective function - often the *evidence lower bound (ELBO)* - and optimizing it using gradient-based methods. Variational inference tends to be faster and more scalable than sampling, though it trades off exactness for speed by approximating the posterior rather than sampling from it directly [BKM17].




Figure 3.1: Comparison of exact posterior to inference methods

## 3.3 Cyber-Physical Systems: Definition and Examples

Cyber-Physical System (CPS) is an umbrella term broadly referring to any systems which tightly integrate both computational logic (cyber) and physical components, allowing interplay through designated interfaces. [Lees5] The two parts of a CPS have to be seamlessly interconnected enabling efficient communication and consistent information relying?? [MS20], as both sides heavily depend on one another. The cyber-physical relationship can be generally describes by three roles *controller/agent*, *sensors* and *actuators*, which are connected by an underlying network. Controllers, representing the 'cyber', are occupied with computation and relying fitting instructions based on te observational data provided by the physical sensors. Sensor and Actuators embody the 'physical' and are the ones interacting with the environments in which the system resides. A sensors jobs is to gather intel from its surroundings communicating it in order for the agent to make correct decisions and predictions and to monitor the system. Without sufficient data, the controller is left blind and possibly unable to perform as intended. While sensors passively interact with the environment by observing, actuators do so actively by altering the systems state — carrying out tasks instructed by the controller. ??Looping behaviour??

Controllers may range from simple reactive routines to advanced inference models grounded in probabilistic programming, or even intelligent agents capable of autonomous decision-making. Actuators, in turn, translating these decisions into physical action can be represented by basic devices such as motors that move parts, valves that regulate flow, or relays that control circuits. They can also act as more complex assemblies — robotic joints, programmable hydraulic systems, or smart actuators that adjust based on local conditions. At the foundation, sensors continuously monitor the environment, capturing data through modalities such as motion, orientation, or spatial mapping. Together, these elements form an integrated feedback system, enabling CPS to respond intelligently and flexibly to their physical context, regardless of application domain.




Figure 3.2: General architecture of common CPS

CP-Systems are inherently multidisciplinary structures intersecting engineering domains (mechanical, electrical, civil, etc.) with computer science. Because these domains employ differing models for ... , effectively combining them is non-trivial. Challenges arise during development and operation of CPSs for which finding solutions is imperative. [MS20] First and foremost is the concern of **reliability and safety**. Occurrences of failures can have dire real-world consequences with catastrophic end results. This means systems must handle unexpected conditions gracefully and robustly. But ensuring safety through extensive verification and validation is complicated by the combination of continuous physical dynamics and discrete decision logics. Another challenge is **real-time performance**, as computing elements must keep pace with the physical processes, which usually require low-latency responsiveness including the network. Since CPSs are used in infrastructure and other critical sites that may be targets of cyber attacks, guarantying a high standard of **Security** is essential. Additionally, sensors are not completely reliable — possibly introducing noise to measurements — and environmental conditions can vary drastically leading to **uncertainty**, which a CPS has to cope with to be resilient. [Thro2]

### 3.4 Motivation for PPLs in CPS Contexts

??enumerate the four paragraphs into 2x2??

Probabilistic programming has multiple fitting applications within CPSs, especially where reasoning under uncertainty is essential. Generally, probabilistic programming allows controllers to plan even in unpredictable environments and settings.

When measurements contain noise or are incorrect, PPL models can still be applied to extract useful information. Instead of treating the raw reading of a sensor as the “truth,” the readings are considered samples for a *sensor’s model*. Not only is it possible to reduce the noise in the measurements by calculating the posterior, but those noise characteristics are also used to flesh out the predictive model itself. This enables PPLs to adaptively adjust noise models online when, for example, sensors deteriorate over time. Additionally — with most CPSs employing multiple sensors — a PPL naturally supports *sensor fusion*, [cpsSensorFusion] combining data from heterogeneous sources to produce a coherent estimate. By defining each sensor’s observation as a conditional distribution, a probabilistic program weights each reading according to its uncertainty: more reliable sensors contribute more heavily, while inconsistent or faulty readings are down-weighted or flagged. Inference then yields a joint posterior over the true state, allowing the system to reconcile disagreements between sensors, detect potential sensor failures by inferring latent bias variables, and maintain a quantified confidence.

Which leads us into robotics, a quintessential cornerstone of CPS. Instead of working on raw observational data and making assumptions in incomplete cases, robots can adopt a *belief* based on an observational model [Thro2]. The observational model couples motion dynamics (how actions translate into state transitions) with existing sensor models, providing a full posterior distribution that the robot can use to imply different levels of confidence for its actions. Furthermore, robots often operate in changing environments or with components that degrade over time. Within a PPL, unknown parameters, such as friction coefficients or sensor calibration

offsets, can be treated as latent variables and perform Bayesian learning to infer them from incoming data. The same inference mechanism used for state estimation thus also adapts the model online, continuously refining both motion dynamics and sensor mappings.

A PPL expresses system dynamics probabilistically so that sampling yields a distribution of possible futures. [Hei+18] By enforcing “chance constraints” (e.g., “limit the probability of exceeding a safety threshold to 1%”), inference or Monte Carlo sampling identifies control inputs that optimize expected performance while respecting risk. When the current state is uncertain, planning occurs in belief space — each action triggers a Bayesian update of the state distribution, and planners choose sequences that minimize expected cost given that uncertainty. Simultaneously, online Bayesian learning refines model parameters (like friction or drag) over time, tightening future predictions; as uncertainty shrinks, control can push closer to limits, whereas rising uncertainty leads to more conservative actions.

Recent research is looking at ways to utilize PPLs to gather diagnostics of a CPS. [PB24] They authors propose a way to infer potential error causes, that might be overlooked otherwise, due to the . This of of great interest to system operators for deciding if either parts have to be replaced or only have to be recelebrated. Because measured data on system deviations often reflects complex intertwining causes and *closed-loop corrections*, identifying root causes becomes a matter of inference. The proposed two-step approach first builds a generative model (based on control software and expert insights) that simulates observations from hypothesized error causes. This simulator is then recast as a probabilistic program, allowing Bayesian inference to estimate latent causes (with confidence intervals) from actual measurements.

## The *Miniprob* DSL

4.1 Domain and Purpose

4.2 Syntax Overview

4.3 Semantics and Example Models

# Language Servers, Parsers & Type Checking

<https://code.visualstudio.com/api/language-extensions/overview>

- 5.1 Language Server Protocol (LSP) Basics
- 5.2 Parser and Syntax Checking
- 5.3 Type Checking Functionality
- 5.4 Overview of Langium



## Part II

# Design & Implementation

# Requirements & Technology Survey

- 6.1 Functional Requirements
- 6.2 Non-functional Requirements
- 6.3 Alternative Technologies
- 6.4 Justification for Langium & VS Code

# Architectural Design

7.1 High-Level Architecture Diagram

7.2 Module Decomposition

7.3 Data Flow and Control Flow

# Implementation Details

8.1 Langium Grammar Definition for *Miniprob*

8.2 Semantic Checks and Type System

8.3 VS Code Extension Points

# Testing & Validation

9.1 Unit Tests

9.2 Integration Tests

9.3 Case Studies & Examples

## Part III

# Evaluation & Discussion

## Performance & Usability

10.1 Parsing Speed

10.2 Editor Responsiveness

10.3 User Feedback

# Discussion

## II.1 Meeting the Requirements



# Bibliography

- [Aho+06] Alfred V. Aho et al. *Compilers: Principles, Techniques, and Tools*. 2nd. Reading, Massachusetts: Addison–Wesley, 2006. ISBN: 978-0321486813.
- [Bac+60] J. W. Backus et al. “Report on the algorithmic language ALGOL 60”. In: *Commun. ACM* 3.5 (May 1960), pp. 299–311. ISSN: 0001-0782. DOI: [10 . 1145 / 367236 . 367262](https://doi.org/10.1145/367236.367262). URL: <https://doi.org/10.1145/367236.367262>.
- [Biso6] Christopher M. Bishop. *Pattern Recognition and Machine Learning*. Springer, 2006.
- [BKM17] David M Blei, Alp Kucukelbir, and Jon D McAuliffe. “Variational Inference: A Review for Statisticians”. In: *Journal of the American Statistical Association* 112.518 (2017), pp. 859–877.
- [DKV00] Arie Deursen, Paul Klint, and Joost Visser. “Domain-Specific Languages: An Annotated Bibliography”. In: *SIGPLAN Notices* 35 (Jan. 2000), pp. 26–36.
- [Goo13] Noah D. Goodman. “The principles and practice of probabilistic programming”. In: *SIGPLAN Not.* 48.1 (Jan. 2013), pp. 399–402. DOI: [10 . 1145 / 2480359 . 2429117](https://doi.org/10.1145/2480359.2429117). URL: <https://doi.org/10.1145/2480359.2429117>.
- [GS14] Noah D. Goodman and Andreas Stuhlmüller. *The Design and Implementation of Probabilistic Programming Languages*. <http://dippl.org>. Accessed: 2024-01-01. 2014.
- [Hei+18] Tor Aksel N. Heirung et al. “Stochastic model predictive control — how does it work?”. In: *Computers and Chemical Engineering* 114 (2018). FOCAPO/CPC 2017, pp. 158–170. DOI: [10 . 1016 / j . compchemeng . 2017 . 10 . 026](https://doi.org/10.1016/j.compchemeng.2017.10.026). URL: <https://doi.org/10.1016/j.compchemeng.2017.10.026>.
- [Jia+09] Tao Jiang et al. “Formal Grammars and Languages”. In: (Nov. 2009). DOI: [10 . 1201 / 9781584888239 - c20](https://doi.org/10.1201/9781584888239-c20).
- [Jina] Pete Jinks. *BNF/EBNF Variants*. <https://www.cs.man.ac.uk/~pjj/bnf/ebnf.html>. Last modified 12 March 2004. Accessed 26 May 2025.
- [Jinb] Pete Jinks. *Notations for Context-Free Grammars: BNF, Syntax Diagrams, EBNF*. <https://www.cs.man.ac.uk/~pjj/bnf/bnf.html>. Last modified 12 March 2004. Accessed 26 May 2025.
- [Lee15] Edward A. Lee. “The past, present and future of cyber-physical systems: a focus on models”. In: *Sensors (Basel)* 15.3 (Feb. 2015), pp. 4837–4869. DOI: [10 . 3390 / s150304837](https://doi.org/10.3390/s150304837).
- [Mee+21] Jan-Willem van de Meent et al. *An Introduction to Probabilistic Programming*. 2021. DOI: [10 . 48550 / arXiv . 1809 . 10756](https://doi.org/10.48550/arXiv.1809.10756). URL: <https://arxiv.org/abs/1809.10756>.
- [MS20] Abha Mahalwar and Rishabh Sharma. “Cyber-Physical Systems: Challenges and Future Directions”. In: *Turkish Journal of Computer and Mathematics Education (TURCOMAT)* 11 (Dec. 2020), pp. 2865–2870. DOI: [10 . 61841 / turcomat . v11i3 . 14651](https://doi.org/10.61841/turcomat.v11i3.14651).
- [PB24] A. Piedrafita and L. Barbini. “PHM Society European Conference”. In: vol. 8. 1. June 2024, p. 7. DOI: [10 . 36001 / phme . 2024 . v8i1 . 4055](https://doi.org/10.36001/phme.2024.v8i1.4055).
- [SK95] Kenneth Slonneger and Barry L. Kurtz. “Specifying Syntax”. In: *Formal Syntax and Semantics of Programming Languages: A Laboratory Based Approach*. Reading, Massachusetts: Addison–Wesley, 1995. Chap. 1, pp. 1–30. ISBN: 0-201-65697-3.

- [Thro2] Sebastian Thrun. “Probabilistic robotics”. In: *Communications of the ACM* 45.3 (Mar. 2002), pp. 52–57. DOI: [10.1145/504729.504754](https://doi.org/10.1145/504729.504754). URL: <https://doi.org/10.1145/504729.504754>.
- [Van12] Peter Van Roy. “Programming Paradigms for Dummies: What Every Programmer Should Know”. In: (Apr. 2012).
- [WGM08] Hui Wu, Jeff Gray, and Marjan Mernik. “Grammar-driven generation of domain-specific language debuggers”. In: *Softw., Pract. Exper.* 38 (Aug. 2008), pp. 1073–1103. DOI: [10.1002/spe.863](https://doi.org/10.1002/spe.863).

## Expanded Code Listings

## Test Data